

EDUCATION

McGill University <i>B.A. Computer Science; Minor in Economics</i> - GPA: 3.8/4.0, Graduated with Distinction (top 25%). - Graduate-level coursework: Applied Machine Learning, Computer Networks, Compiler Design. - Girls Who Code (McGill) Vice President; co-developed a Python curriculum; taught weekly advanced-level high school classes.	Montreal, Canada 2024 – 2026
IE University <i>B.Sc. International Relations</i> - GPA: 9.0/10.0, Dean’s List (top 10%); transferred to McGill after completing the second year. - Honours: Comparative Politics, International Law, Quantitative Methods, International Political Economy, Foreign Aid. - Research assistant to Karen Nershi: studied statistical patterns in state-backed ransomware activity during electoral cycles.	Madrid, Spain 2022 – 2024

PUBLICATIONS

Ceccatelli, V., Jeon, Y., & Adelani, D. I. (2026). SpeechJBB: Probing Safety Alignment and Comprehension in Large Audio Language Models under Code-Switched Speech. *Proceedings of EMNLP 2026, Main Conference*.

Jeon, Y., Maltais, M., **Ceccatelli, V.**, Ma, M., & Adelani, D. I. (2026). VoxSumm: A Multilingual Corpus of Long-Form Spoken News for Joint Summarization and Translation. *Submitted to EACL 2027*.

Ceccatelli, V. & Diop, F. (2026). Securing the Digital Frontier: U.S.–Sub-Saharan Africa Cybersecurity Partnerships in the Shadow of China’s Digital Silk Road. *Andalus Committee Policy Paper*.

Research assistance: Nershi, K. (2025). State–cybercriminal alignment in the ransomware ecosystem. *Journal of Cybersecurity*.

Ceccatelli, V. (2024). Sky Diplomacy: The Geopolitical Impact of the Proliferation of Iranian-Russian Military Drone Trade on Global Alliances and Security. *IE International Policy Review*.

RESEARCH EXPERIENCE

UCL S2Lab, Prof. Lorenzo Cavallaro <i>AI Security Researcher</i> - Co-developed and validated a taxonomy of internal LLM uncertainty signals, introducing trajectory-based probes that improved uncertainty estimation across four model families while reducing inference cost relative to sampling-based methods. - Built a framework to trace binding, data flow, control dependence, and security taint flow in code models, benchmarking ~2,000 synthetic and real Python programs with causal interventions and leakage-resistant controls. - Verifying causal use of semantic representations via activation patching, and stress-testing robustness under semantics-preserving transformations across 5 obfuscation levels, enabling detection of semantic degradation before it surfaces in outputs.	London, UK (Remote) June 2026 – Present
MILA Quebec AI Institute, Prof. David Adelani <i>AI Safety Researcher</i> - Led SpeechJBB, the first audio-based code-switching jailbreak dataset for multilingual speech-safety evaluation, exposing weaknesses in SOTA LALM safety alignment; validated findings against multilingual reasoning and ASR capabilities. - Identified that non-English code-switching increased mean jailbreak success by 28% and reduced refusal by 14.4% relative to monolingual speech; introduced a further pseudo-word obfuscation attack increasing jailbreak success by 34%. - Created the dataset underpinning Voxsumm, a multilingual corpus for joint summarization and translation of long-form spoken news in low-resource languages, developed in collaboration with Google DeepMind.	Montreal, Canada January 2026 – August 2026
McGill Prometheus Lab, Prof. Joseph Vybihal <i>Computer Vision Researcher</i> - Devised a ground-segmentation and motion-planning system in C++ (U-Net, MobileNetV2), deployed across images, video, and live webcam for autonomous indoor navigation, reaching 0.91 validation IoU on a custom-labelled dataset. - Selected the deployment checkpoint via safety-driven evaluation, prioritising small-obstacle detection and calibrated confidence on ambiguous surfaces (shadows, reflections) over peak IoU to ensure navigational safety.	Montreal, Canada May 2025 – August 2025

PROFESSIONAL EXPERIENCE

WIIT - The Premium Cloud <i>AI Security Intern</i> - Collaborating with WIIT AI Engineers to build an agentic purple-teaming loop that continuously red-teams AI agents and autonomously hardens them, integrating HikmaAI’s enforcement layer as the runtime control. - Spearheading automated attacker with defined taxonomy, including prompt injection, tool-parameter manipulation and goal hijacking, as well as safety pillars such as bias and misinformation; producing an evaluation harness to measure success. - Quantifying detection, bypass and false-positive rates across [X] attack vectors and reducing agent bypass rate by [Y] %.	Düsseldorf, Germany July 2026 – Present
Wavestone <i>Digital Transformation Intern</i> - Designed an RL system, including environment, state & action spaces, and reward function, to automate data cleaning tasks; reduced records requiring manual review by 29% while preserving 96% correction precision. - Improved the overall validity, accuracy and completeness of financial datasets; deployed across German actuarial department.	Munich, Germany June 2024 – July 2024

ADDITIONAL SKILLS

- Technical Skills: Python, C/C++, Java, OCaml, Assembly, Bash, SQL, R, Stata, MATLAB; PyTorch, Hugging Face Transformers, scikit-learn, NumPy, pandas, OpenCV; Git, Docker, Linux; CUDA, SGE/SLURM, LLM APIs; program analysis.
- Languages: Italian (C2). English (C2). German (C2). Spanish (C1). French (C1).